

# CONVOLUTIONAL NEURAL NETWORK LANGUAGE MODELS

A DISSERTATION  
SUBMITTED TO THE DEPARTMENT OF COMPUTER SCIENCE  
AND THE COMMITTEE ON GRADUATE STUDIES  
OF UNIVERSITY OF TRENTO  
IN PARTIAL FULFILLMENT OF THE REQUIREMENTS  
FOR THE DEGREE OF  
MASTER OF SCIENCE

Ngoc-Quan Pham  
September 2016

© Copyright by Ngoc-Quan Pham 2016  
All Rights Reserved

I certify that I have read this dissertation and that, in my opinion, it is fully adequate in scope and quality as a dissertation for the degree of Master of Science.

---

(Marco Baroni) Principal Adviser

I certify that I have read this dissertation and that, in my opinion, it is fully adequate in scope and quality as a dissertation for the degree of Master of Science.

---

(Gemma Boleda)

I certify that I have read this dissertation and that, in my opinion, it is fully adequate in scope and quality as a dissertation for the degree of Master of Science.

---

(German Kruzhevsky)

---

# **Acknowledgements**

# Contents

<b>Acknowledgements</b>	<b>iv</b>
<b>1 Introduction</b>	<b>1</b>
<b>2 Background</b>	<b>5</b>
2.1 Statistical Language Modeling . . . . .	5
2.1.1 Problem overview . . . . .	5
2.1.2 Evaluation Metrics . . . . .	6
2.1.3 $N$ -gram models . . . . .	7
2.2 Neural Network Language Models . . . . .	9
2.2.1 Feed-forward neural language models . . . . .	10
2.2.2 Training method . . . . .	14
2.2.3 Optimisation Process . . . . .	14
2.2.4 Recurrent neural language models . . . . .	17
2.3 Convolutional Neural Networks . . . . .	27
2.3.1 The convolutional layer . . . . .	27
2.3.2 Hyper parameters for Convolutional Neural Network . . . . .	29
<b>3 Convolutional Neural Language Models</b>	<b>31</b>
3.1 Base FFLM . . . . .	32
3.2 CNN and variants . . . . .	33
<b>4 Experiments</b>	<b>37</b>
4.1 Experiment details . . . . .	37

4.2	Results . . . . .	38
<b>5</b>	<b>Model Analysis</b>	<b>41</b>
<b>6</b>	<b>Related Works</b>	<b>45</b>
<b>7</b>	<b>Conclusion</b>	<b>47</b>
<b>A</b>	<b>Publications by Author</b>	<b>49</b>

# List of Tables

2.1	An example of perplexity computation for the sentence “The relationship between Obama and Netanyahu is not exactly friendly” . . . . .	6
2.2	Notations for neural network layers. . . . .	12
2.3	Notations for recurrent neural network layers. . . . .	20
4.1	Results on Penn Treebank and Europarl-NC. Figure of merit is perplexity (lower is better). Legend: $k$ : embedding size; $ker$ : kernel size; $val$ : results on validation data; $test$ : results on test data; $\#p$ : number of parameters; $hid$ : size of hidden layers; $\#l$ : number of layers. . . . .	40

# List of Figures

2.1	Neural Language Model by Bengio et al [3]. . . . .	11
2.2	A simple RNNLM [38] predicting the sequence “The cat is eating fish bone”. . . . .	19
2.3	Simple illustration for convolution. The input is a 2D image, the output is obtained by sliding the kernel through the image. . . . .	28
3.1	Overview of baseline FFLM. . . . .	33
3.2	Convolutional layer on top of the context matrix. . . . .	35
3.3	Combining kernels with different sizes. We concatenate the outputs of 2 convolutional blocks with kernel size of 5 and 3 respectively. . . . .	36
5.1	Some example phrases that have highest activations for 8 example kernels (each box), extracted from the validation set of the Penn Treebank. Model trained with 256 kernels for 256-dimension word vectors. . . . .	41
5.2	The distribution of positive weights over context positions, where 1 is the position closest to the predicted word. . . . .	42
5.3	Perplexity change over position, by incrementally revealing the Mapping’s weights corresponding to each position. . . . .	44



# Chapter 1

## Introduction

Exploiting the information from textual data (corpora) for artificial intelligent systems is one of the most important targets of Natural Language Processing (NLP). Applications that can benefit from this textual knowledge include Machine Translation, Information Extraction, Search Engine, Automatic Speech Recognition or Text to Speech ...

The language modeling problem was formulated in order to successfully extract the knowledge about languages from the corpora. The main purpose of a language model is to model the linguistic regularities of language by capturing the morphological, syntactical and semantic properties of a given language. Such a model should ideally be able to assess the likelihood of an arbitrary sentence in a particular context. In the last decades, the language modeling problem has been approached with statistical methods. As a standard, the language models are trained so as to predict the upcoming word given its previous context, which are the words appearing before the word to be predicted. The word and its context is usually referred as an  **$n$ -gram**, which is a contiguous sequence of  $n$  tokens (words in textual data), consisting of one word to be predicted and  $n - 1$  words in the given context. For example, the word combination “The dog runs” is an  $n$ -gram with order of 3, and the language models aim at predicting the word “runs” given the context “the dog”.

The conventional approach to this problem relies on the statistical point of view, which is assigning probabilities for sentences by counting the number of occurrences of  $n$ -grams in the corpora to estimate the probabilities. Problems arise when the model has to estimate the probabilities of word sequences that are rare in the corpus, or do not appear at all, but

still make sense in human languages. There are techniques that estimate the distribution of rare  $n$ -grams based on the seen ones (also referred as smoothing techniques) or using the statistics of lower order  $n$ -grams (back-off techniques) [9, 14, 22, 30, 41, 61].

The main disadvantage of these models is that each individual word is treated independently of the others, count-based  $n$ -gram models fail to capture semantic relations between words. As a result, those models suffer from two weaknesses. First, they are limited by a small context size, since the number of possible  $n$ -grams grows exponentially with the  $n$ -gram order. For example, if the language vocabulary contains 10000 words then there are  $10^{24}$  possible 6-grams to be estimated, which cannot be covered by any corpora as a consequence of the Zipf's law [29]. Second, it is not able to infer the estimation of an unknown word string based on similar sequences. For example, if a good language model has seen the sentence "The cat runs", then it should be able to tell that the unseen sentence "A dog runs" is valid, given the semantic similarity between the articles "a" and "the", and the pets "cat" and "dog".

In contrast neural network language models [4] became a solution for the problems of  $n$ -gram models. The success of neural network language models is achieved based on two key ideas. First, the models learn the *distributed representation* of words, which express the word meaning with low dimensional vectors which can represent the similarity structure between words. Second, the statistical language model is viewed as a machine learning classification problem, in which the model tries to *discriminate* the predicted word from other words in the vocabulary. The models, initially proposed by Bengio et al. [4] and then enhanced by Hai Son et al. [19], Mikolov et al. [38], Schwenk [49], Vaswani et al. [59] have been shown to outperform the count-based  $n$ -gram models.

Neural networks (or Artificial neural networks) are a class of machine learning models that is inspired by biological neural network. They are well-known to be flexible and powerful in terms of learning representations in various problems, including Language Processing. Among the commonly used architectures of neural networks, there are two different structures applied successfully in language modeling: feed-forward networks and recurrent networks, making different architectural decisions.

Recurrent networks [38, 39] treat language modeling as a sequence prediction problem by taking one word token at a time together with a hidden "memory" mechanism to produce

a predicted word and update the memory for next time step. Importantly, they are strong enough to represent any  $n$ -gram at arbitrary size. Contrarily, feed-forward models take as input the last  $n - 1$  words in the  $n$ -gram, where  $n$  has to be a fixed window size, and use them jointly to predict the upcoming word.

Despite the fact that the state-of-the-art language models consist of a shared vector space between words, they still treat the context as a sequence of discrete symbols. The feed-forward network simply concatenate the word vectors and encodes the context into the hidden layer, while the recurrent models take one word as a time step, whose weakness is that a particular time step does not have any information about future steps. In fact, it is useful to take into account the collocations and multi-word expressions in the context. For example, the English word “get” can acquire different meanings when paired with different prepositions, which the word embeddings hardly can represent. Such word combinations are often identified by statistical methods [48], which requires an intensive scanning through the data to find frequent combinations. In the context of the neural networks, it is intuitive to use a special neural network architecture that can search the context for patterns and combinations.

In the literature of machine learning and computer vision, Convolutional Neural Networks (CNNs) are the family of neural network models that feature a type of layer known as the convolutional layer. The main idea of this layer is to look at each position of the input (at any dimension) with a fixed size window and find local features - distinctive attribute that are useful for the learning task. The network has been applied successfully in image processing and speech processing [18] but they have received less attention in Natural Language Processing. Mainly, they have been somewhat explored in *static classification tasks* where the model is provided with a full linguistic unit as input (e.g. a sentence) and classes are treated as independent of each other. Examples of this are sentence or document classification for tasks such as Sentiment Analysis or Topic Categorization [27, 28], sentence matching [25], and relation extraction [42]. However, their application to *sequential prediction tasks*, where the input is construed to be part of a sequence (for example, language modeling or POS tagging), has been rather limited (with exceptions, such as Collobert et al. [11]).

The language modeling in the work of Collobert et al. [11] utilised the convolutional

layers to extract features for multiple natural language processing tasks including language modeling. However, due to the computational limit, the model was altered to focus on learning vectorised word representation, leaving the language modeling quality a question mark. In this thesis, we investigate the **feed-forward** neural language models with the temporal convolutional layer inspired by the work of [11]. Starting by analysing the existing neural language models including feed-forward and recurrent ones, we investigate the convolutional models with some reasonable changes in terms of architecture to verify if the convolutional neural network deliver any improvement over the feed-forward counterpart. For that purpose, we initially train a carefully tuned feed-forward neural language model, which yield competitive performance in various corpora. Subsequently, we enhance the baseline with an additional convolutional layer that convolves over the word embeddings, and we can effortlessly limit the customization space for network architecture during training. The experimental results indicate that using convolutional layers can improve 15-20% perplexity compared to the solid baseline, and the CNN model can perform on par with the state-of-the-art models. Our analysis also shows that the convolutional layer independently learned to focus on different linguistic patterns, and the model can take into account context words at very far from the target.

# Chapter 2

## Background

### 2.1 Statistical Language Modeling

#### 2.1.1 Problem overview

As we introduced before, the Statistical approach in Language Modeling aims at measuring the fluency of all possible word strings, which are considered as a stochastic process, with probabilities. The sequence length can be arbitrary, while the words are taken from a limited vocabulary. A trained language model should be able to show that the likelihood of "the end of our world" is much higher than "tea end of our word", because the latter string is much less likely to be found in available English text. Concretely, let us denote  $W_1^L$  to be a word sequence with length  $L$  which is composed by the constituent words  $w_1, w_2, \dots, w_L$ , then a statistical language model aims at predicting the probability distribution of all possible sequences  $W_1^L$  in a particular language:

$$P(W_1^L) = P(w_1 w_2 \dots w_L) \quad (2.1)$$

Since the direct estimation for that probability distribution is intractable, the probability of a sentence  $P(w_1^L)$  is factorised using the chain rule:

$$P(W_1^L) = P(w_1 | <s>) \prod_{i=2}^L P(w_i | <s> W_2^{L-1}) = \prod_{i=1}^L P(w_i | H_i) \quad (2.2)$$

Word to be predict	Context	Probability (example)
The	<s>	0.1
relationship	<s>The	0.002
between	<s>The relationship	0.05
Obama	<s>...relationship between	0.00001
and	<s>...between Obama	0.2
Netanyahu	<s>...Obama and	0.00001
is	<s>...and Netanyahu	0.05
not	<s>...Netanyahu is	0.02
exactly	<s>...is not	0.001
friendly	<s>...not exactly	0.0007
.	<s>...exactly friendly	0.2
Perplexity		210

Table 2.1: An example of perplexity computation for the sentence “The relationship between Obama and Netanyahu is not exactly friendly”.

in which,  $\langle s \rangle$  is used to denote beginning token of the sentence/string. The history  $H_i$  represents the string before the current word  $w_i$ . Instead of directly modeling the original distribution, the target probability is factorised into constituent *conditional* probabilities, which are more practical to estimate.

### 2.1.2 Evaluation Metrics

The quality of statistical language models can be evaluated by the capability to predict a new corpus, which is defined by using *Perplexity*. Let us assume that we have a trained language model  $M$  and a corpus  $D$  containing  $L$  words, which the language model has not observed during the training process. The quality of model  $M$  is evaluated by using it to predict the distribution of the corpus  $D$ . The resulted perplexity (PPL) is then obtained by estimating the probability of all words given their context in the corpus  $D$ , we use  $P_M$  to denote the probability distribution produced by model  $M$ . An example of PPL computation is illustrated on Table 2.1.2.

$$PPL(D) = \exp\left(\frac{\sum_{l=1}^L -\ln P_M(w_l|H_l)}{L}\right) \quad (2.3)$$

The idea of perplexity is that, minimising perplexity is corresponding to maximising

the ability to predict every word in the data. A low perplexity value corresponds to the fact that the language model is able to fit better the data, since the distribution of the model is closer to the unknown distribution of the test data. One property of perplexity is that, it is correlated to the average number of guessing needed to predict all words in the data. Therefore, a random distribution or the “worst” model (usually the model with parameters initialised) often gives the perplexity very close to the vocabulary size  $V$ . This property is often used to check the implementation of the model and perplexity.

In practice, the quality of language models can also be evaluated through their impacts on other applications such as Automatic Speech Recognition (ASR) or Statistical Machine Translation (SMT), by reducing the errors in the output of such systems. For example, two strings “Tea end of our word” and “The end of our world” may have similar speech signals and the recognizer has to rely on a good language model to distinguish them.

The main advantage of perplexity is that it is fast to perform and independent to other complex systems. Therefore, in this thesis, PPL would be the only measurement used in the experiments. However, it is important to note that an improvement in terms of perplexity does not always result in the application improvement. For example, the improvement is required to be at least 10% to be noteworthy for an ASR system [46]. Also, there are some language models that aim at ranking the words in the vocabulary by using unnormalised distributions (the sum of the probabilities does not equal to 1). In that case, PPL is also unusable.

### 2.1.3 $N$ -gram models

Going back to equation 2.2, the target of statistical language models is to estimate the conditional probabilities  $P(w_i|H_i)$  with  $H_i$  being the history context of  $w_i$  ( $i$  is an arbitrary index in the corpus). A common approach is to limit the length of the history  $H_i$  by the Markov assumption: the meaning of the word  $w_i$  only depends on  $n - 1$  words before, thus limiting the history for all words in the data into  $n$ -grams:

$$P(w_i|H_i) \approx P(w_i|W_{i-n+1}^{i-1}) \quad (2.4)$$

With  $W_j^k$  denotes the word sequence from index  $j$  to index  $k$ . It is natural to use  $n - 1$  to

denote the order of language models because it is the order of the Markov model reflected in Equation 2.4, however in the literature, authors always use  $n$  to denote the order of  $n$ -grams. Therefore, in order to be clear, an  $n$ -gram contains one word to be predicted, and  $n - 1$  words in the context.

Count-based models estimate the probability of each  $n$ -gram based on simple counting with maximum likelihood estimation. Let  $c(W_j^i)$  denote the number of times that this word sequence occurs, the estimated distribution can be derived as follows:

$$P(w_i|W_{i-n+1}^{i-1}) = \frac{c(W_{i-n+1}^i)}{c(W_{i-n+1}^{i-1})} \quad (2.5)$$

The method is unreliable since it struggles at estimating the distribution for rare and unseen  $n$ -grams (the nominator is 0), many of which actually make sense in natural language. Many techniques have been proposed to overcome this weakness, from re-distributing the probabilities of the frequent  $n$ -grams to the less frequent ones (smoothing techniques) or deriving the distributions of the rare  $n$ -grams from the lower order  $n$ -grams (interpolation and back-off techniques) [9, 14, 22, 30, 41, 61]. However, even with complicated smoothing techniques, the main weaknesses of  $n$ -gram language models are still exposed as follows:

First, each word in the vocabulary is treated as a totally discrete random variable without any linguistic feature associated. The model relies on statistical occurrences and ignores morphological, syntactic and semantic relationship, by which the lexicon is formulated. There are several attempts to incorporate the word similarity in  $n$ -gram based language models. Notably, the class-based language models [7, 43] introduced word clustering and assumed that the distribution of unseen or rare words can be achieved by using the richer statistics from the corresponding class, which is less sparse than the word itself. Also, the structured language models [8, 15] tries to filter out irrelevant context words and focuses on important counterparts by using parse trees, which compensates the lack of syntactic information in  $n$ -gram models. Despite such efforts, the language modeling results were still unreliable compared to the Knesey-Ney smoothing technique. However, those works in the literature also suggests that the syntactic and semantic properties of words need to be automatically learned from the data. Second,  $n$ -gram language models struggle to model



a long range dependency between the predicted word and the context. Due to the fact that each word in the vocabulary is a separated random variable, the number of parameters to be estimated (statistics of  $n$ -grams) grow exponentially with the size of context. The *curse of dimensionality* refers to the fact that one needs more amount of training data in order to reliably estimate the model when the number of learned parameters increases. For example, if the vocabulary size is 10000, the total number of  $n$ -grams for  $n = 6$  is  $10^{16}$  theoretically, which is also the number of parameters to be estimated accordingly. Also, the Zipf's law [29] indicates that only a small subset of the vocabulary accounts for most occurrences in the training data, thus it is almost impossible to have a training data that covers all possible  $n$ -grams.

The neural network language model, as a consequence, is investigated in order to tackle both problems that traditional count-based models cannot solve.

## 2.2 Neural Network Language Models

In this section, we describe the neural language models (also known as continuous space language models), which are designed to fight the curse of dimensionality in the conventional approaches, by utilising two properties as follows

- Words are no longer discrete variables without any relationship to each other. Alternatively, they are represented as real-valued vectors in a continuous space, where similar words are neighbors in the space, such approach is referred as *Distributed word representation* or *Word embedding* [3, 38]. With such representation, each  $n$ -gram is a combination of the word vectors and  $n$ -grams with similar words share the same word vectors. In contrary to the count-based approach, if we increase the  $n$ -gram order then thanks to the sharing property, we just need to adjust the size of the context vector, which can grow linearly with the  $n$ -gram order (in the case of feed-forward neural networks) or remains the same (in the case of recurrent neural networks).
- When the  $n$ -gram context  $H_i$  is the combination of the word vectors, the conditional

probability distribution  $P(w_i|H_i)$  of each  $n$ -gram is then expressed as a parameterised smooth function of the context vector. In machine learning, neural networks are a class of models that can be used to approximate functions. As a result, Bengio et al. [3] proposed to incorporate the continuous word space with a neural network architecture in order to learn the smooth function. The training methods of neural network allow us to **jointly learn** the word space and the parameters of that probability function.

In the following, we provide the detailed description of two most successful architectures used for neural language models: the feed-forward models and the recurrent models.

### 2.2.1 Feed-forward neural language models

The standard architecture of a neural language model [3] is illustrated in Figure 2.1. The model takes the context words  $H_i$  as inputs and outputs the conditional probability distribution over *all* words in the vocabulary which expresses  $P(w_i|H_i)$ . It consists of three basic components: the input layer, the hidden layer and the output layer. The notations are covered in Table 2.2.1. In a standard feed-forward neural network, each layer is a real-valued **vector**, while the (learnable) weights (or parameters) are real-valued **matrices** connecting the layers together. Also, in the standard architecture we also consider networks with only one hidden layer for ease of understanding. In practice, it is possible (and beneficial) to extend the network with multiple hidden layers.

**Input layer** The input layer constructs the embeddings of the words in the context, by using a shared word space and mapping each word in the context to a real-valued vector. Concretely, each word  $w_i$  is represented as an 1-of- $V$  coding, which is a long vector with the size of vocabulary  $V$  with all zeros except for the element corresponding to the word's index in the dictionary. Using this form of sparse coding, the word space (also called projection matrix  $R$ ) is a matrix that contains  $V$  rows and each word embedding  $v_i$  corresponds to one row of the matrix. The number of columns of the matrix is the size of the embedding, which is a tunable hyper-parameter. Notably, the values of the word vectors do not depend




Figure 2.1: Neural Language Model by Bengio et al [3].

on the position of the word in the context. The context vector  $i$  is the concatenation of the word vectors.

$$i = \{R^T v_1; R^T v_2; \dots; R^T v_{n-1}\} \quad (2.6)$$

**Hidden layer** In the hidden layer, the input (context vector) is transformed nonlinearly, where each layer activation values are defined by

$$h = f(W^h i + b^h) \quad (2.7)$$

In equation 2.15, the hidden layer  $h$  has the corresponding weights  $W^h$  and  $b^h$ . The input of the hidden layer is the context vector produced from the input (projection) layer. The size of the hidden layers are tunable hyper parameters.  $f$  denotes a nonlinear activation

notation	meaning
$V$	Vocabulary size
$m$	word vector dimension
$n$	order of language model
$w$	a word; its index in the vocabulary
$v_w$	one-hot vector for word $w$ . A vector with size $V$ with all null except the $w^{th}$ index (=1)
$i \in \mathbb{R}^{(n-1)M}$	input vector of the network.
$o$	output layer of the network which expresses the unnormalised probability distribution
$h$	hidden layer of the network
$p$	the output vector of the network which is normalised for the conditional probability distribution
$i_j$	the $j^{th}$ -index of vector $i$ , which is also used for $o$ , $h$ or $p$
$f$	nonlinear function
$H$	hidden layer size
$W^h \in \mathbb{R}^{H \times (n-1)M}$	the weight matrix connecting the input layer with the hidden layer
$b^h \in \mathbb{R}^H$	the bias vector for the hidden layer
$W^o \in \mathbb{R}^{H \times V}$	the weight matrix connecting the hidden layer with the output layer
$b^o \in \mathbb{R}^V$	the bias vector for the output layer
$\mathcal{L}$	The objective (loss) function for learning the network
$dX$	The Jacobian matrix - the matrix of all of the first order partial derivatives of the loss function with respect to the vector/matrix $X$ .

Table 2.2: Notations for neural network layers.

function. Popular choices for the activation function are Tangent Hyperbolic, Sigmoid or ReLU, expressed in equation 2.8.

$$f(x) = \begin{cases} \frac{\exp(x) - \exp(-x)}{\exp(x) + \exp(-x)} & \text{if } f = \text{Tanh} \\ \frac{1}{1 + \exp(-x)} & \text{if } f = \text{Sigmoid} \\ \max(0, x) & \text{if } f = \text{ReLU} \end{cases} \quad (2.8)$$

**Output layer** The final layer of the network produces the probability distribution for all words in the vocabulary, thus having totally  $V$  nodes. Each neuron in the layer is associated to the probability of one word, as shown in Figure 2.1. First, a linear transformation is used to obtain the unnormalised distribution:

$$o = W^o h + b^o \quad (2.9)$$

Notably, the values of  $o$  is unnormalised because each element in  $o$  is associate to the score of each output word given the context vector.  $W^o$  and  $b^o$  are the corresponding weights and biases of the layer. Importantly,  $W^o$  has the same form as the projection matrix at the input layer, since it also learns embedding for each word in the output vocabulary. Subsequently, the true probability distribution is estimated thanks to the *softmax* function:

$$p(w_i|h) = \frac{\exp(o_i)}{\sum_j \exp(o_j)} \quad (2.10)$$

In equation 2.10, the probability of each word  $w_i$  given the encoded context  $h$  is estimated by normalizing all values in  $o$ . Overall, the main tunable hyper-parameters of the network are the order of  $n$ -grams (the number of words in the context, which can range from 6 to 15), the sizes of hidden layers, and the word embedding size. The set of free parameters  $\Theta$  that are iteratively updated by learning from the data includes the projection matrix  $R$ , the weight matrices  $W^h$  at the hidden layers,  $W^o$  at the output layer and the biases  $b^h, b^o$ .

### 2.2.2 Training method

From the machine learning perspective, the neural network language model has transformed the statistical language modeling problem from a generative learning process to a discriminative classification problem. The free parameters of the network are trained by minimising the objective function, which is the log-likelihood  $\mathcal{L}$  of the parameters  $\Theta$  given the training samples. The parameters are updated after each iteration based on some optimization techniques, among which Stochastic Gradient Descent (SGD) is most commonly used in neural language models [19, 38, 63]. SGD and other variants such as Adadelta [64] or RMSProp [58] require the computation of the first order derivatives of the loss function with respect to the parameters, which can be performed efficiently with the back-propagation algorithm [34].

### 2.2.3 Optimisation Process

In this section, we describe the back-propagation flow in the standard feed forward neural language model - the core of the optimisation process. Back-propagation [47] involves using a dynamic programming strategy to compute the derivatives of the loss function with respect to the parameters layer by layer, based on the chain rules. In the standard network, the error derivatives are back-propagated from the output layer to the input (projection) layer.

**Objective Function** The smoothing function that we approximate with the neural network has parameters that can be iteratively tuned in order to **maximise the log-likelihood of the training data** [3]. The objective function is therefore chosen as the Negative Log-Likelihood function, since SGD requires the training objective to be minimised. Assuming we have  $N$  samples in the training data, each of which is an  $n$ -gram, we can compute the loss function over the training data as follows:

$$\mathcal{L} = - \sum_i^N \log P(w_i | w_{i-n+1}^{i-1}) \quad (2.11)$$

The loss function is also in line with the Perplexity in Equation 2.3. For ease of understanding, we denote the derivative of the loss function  $\mathcal{L}$  at *each* sample or mini-batch of samples with respect to a variable  $x \in \Theta$  by  $dx$ .

For each sample  $w_i$  and its context  $H_i$ , we have:

$$\begin{aligned} -\log P(w_i | w_{i-n+1}^{i-1}) &= -\log P(w_i | H_i) \\ &= -\log\left(\frac{\exp(o_w)}{\sum_i \exp(o_i)}\right) \\ &= \log\left(\sum_i \exp(o_i)\right) - o_w \end{aligned} \quad (2.12)$$

Subsequently, we compute the error derivatives  $dx$  given the parameters in each layer using back-propagation:

**Output layer** The derivatives at the output layer:

$$do_i = \begin{cases} 1 - p_i & \text{if } i == w \\ -p_i & \text{otherwise} \end{cases} \quad (2.13)$$

Notably,  $o_i$  denotes the  $i^{th}$  element of the vector  $o$ , which is the unnormalised conditional distribution of the vocabulary given the encoded context  $h$ .

**Hidden layers** As a result, we can compute the derivatives with respect to the parameters and the previous hidden layer  $h$ , based on the original inference from Equation 2.9.

$$\begin{aligned} dW^o &= doh^T \\ db^o &= do \\ dh &= W^{oT} do \end{aligned} \quad (2.14)$$

**Input Layer** The inference equation for the hidden layer from the input layer:

$$h = f(W^h i + b^h) \quad (2.15)$$

which implies that:

$$\begin{aligned}
 d[W^h i + b^h] &= f'(h) * dh \\
 db^h &= d[W^h i + b^h] \\
 dW^h &= d[W^h i + b^h] i^T \\
 di &= W^{hT} d[W^h i + b^h]
 \end{aligned} \tag{2.16}$$

In order to have the derivatives for the activation function  $f$ , we have:

$$f'(x) = \begin{cases} 1 - \text{Tanh}(x)^2 & \text{if } f = \text{Tanh} \\ \text{Sigmoid}(x) - \text{Sigmoid}(x)^2 & \text{if } f = \text{Sigmoid} \\ 1 \text{ when } x > 0 \text{ and } 0 \text{ otherwise} & \text{if } f = \text{ReLU} \end{cases} \tag{2.17}$$

**Parameter Update** After obtaining the derivatives of the loss function with respect to all parameters in the network, we can update the parameters following to Stochastic Gradient Descent. The method is based on the phenomenon that the gradient of a function always points towards the direction of maximal increase at any point. The update rule is as follows with the learning rate parameter  $\alpha > 0$  and an arbitrary parameter  $x$ :

$$x = x - \alpha dx \tag{2.18}$$

The learning rate is also considered as a function of the number of samples trained in the data. From experiments, the learning rate is updated after the model observes a number of training examples with two typical ways. The first way is to exponentially decrease the learning rate after some training samples with a *learning rate decay*, normally an epoch (training all samples in the training data). The second way is to reduce the learning rate based on a validation data. After each epoch, if the perplexity on the validation data is decreased, the learning rate is kept the same, otherwise it is multiplied by the learning rate decay.



**Methods to prevent overfitting** Overfitting is a phenomenon that the model has poor predictive performance, even if the model is well trained on the training data. The possibility of overfitting exists because the criterion used for training the model is different than the criterion used to measure the efficacy of the model. It is very likely that the training data yields a different distribution than the test data, therefore fitting the model on the training data does not guarantee a good prediction performance.

For neural network language models, the main methods used to prevent this phenomenon to happen is to apply **regularisation methods** and **early-stopping** strategy.

- **Regularisation:** The most efficient method is to apply **Dropout technique**, which refers to dropping out units (in most case, hidden units) in neural networks [23, 53]. Concretely, we temporarily set the unit values to 0 based on a random distribution (Bernoulli distribution, for example) during the training phase. In the testing phase, the unit is always present. Dropout is usually applied in the hidden layers of feed-forward neural networks. The choice of which unit to drop in the layer is random which is normally associated with a fixed probability  $p$  independent for each unit. Depending on the network size and the amount of training data, the value of  $p$  is chosen empirically.
- **Early-stopping:** In order to prevent overfitting, we use a validation set during training. The validation data is a separate set with the size similar to the test data. After each epoch (normally a whole scan over the training data), we measure the perplexity of the model on the validation data. If the perplexity does not decrease, the learning rate is reduced in the next epoch. The training process is halted when the learning rate reaches a threshold.

## 2.2.4 Recurrent neural language models

Compared to the statistical  $n$ -gram models, the feed-forward neural language models had created a considerable leap in representation by combining distributed representation of words with a robust classifier to generalise from observed sequences. As can be seen from various works [19, 49], the feed-forward language model significantly outperformed the

traditional count-based models. However, the feed-forward models require a fixed input size, thus still rely on the Markov Assumption which limits the context to a particular number of words, even when the context size can be large. In order to model long sentences, or even paragraphs with long-term dependencies, it is beneficial to investigate in models that can be flexible in terms of input size. For example, if the distance between the open bracket and the closed counterpart is further than the  $n$ -gram input size, the feed forward model may forget to close the brackets after seeing the initial one. Ideally, as the learning process of human is associated with a memory that keeps the current information (such as topics), similar structure should be simulated and integrated in the network.

**Recurrent Neural Networks** RNN [13] are a class of neural networks that can efficiently model sequences by using a dynamic memory structure. While the feed-forward network can only receive one input and compute the corresponding output without any relation with other inputs, the recurrent counterpart takes the input as a *series of time step*  $x_1, x_2, \dots, x_n$  and processes them one by one, taking into account the information stored in the previous steps. Concretely, for each input  $x_i$ , the network updates the hidden memory  $h^i$  based on the previous one  $h^{i-1}$ .

$$h^i = \mathcal{F}(x_i, h^{i-1}) \quad (2.19)$$

We will cover some popular RNN variations in the upcoming sections. In general, the strong point of these models lie in the ability to dynamically model sequences with arbitrary length, which the feed-forward neural networks cannot achieve. The advantage comes with the cost that the recurrent models are generally hard to train, due to the properties of back-propagation. A change in an arbitrary position of the sequence can lead to a change in the objective function, therefore training methods for RNNs typically have to trace back the previous time steps. In other words, the RNNs are equivalent to feed-forward neural networks with many hidden layers that share parameters across each other. Importantly, the model capacity of RNNs do not depend on the length of the sequences, but in the recurrence mechanism - the way the hidden layers are updated.

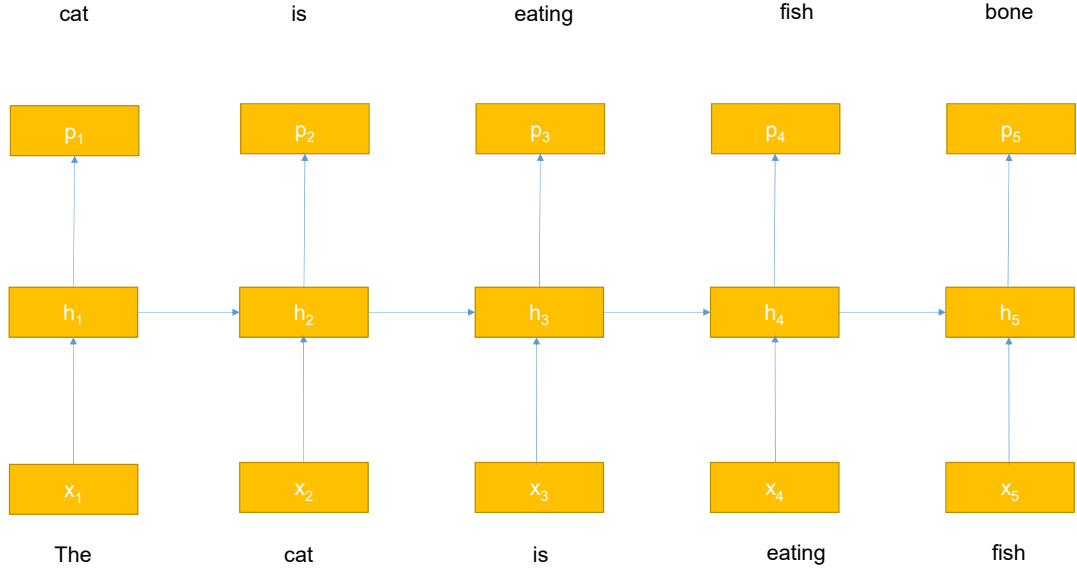


Figure 2.2: A simple RNNLM [38] predicting the sequence “The cat is eating fish bone”.

**Recurrent Language Models** Language Modeling can be viewed as a sequence modeling problem, in which each time step corresponds to one word. We reuse the notations from Table 2.2.1 with another notation for recurrent layers added by Table 2.2.4

The first recurrent language model (RNNLM) [38] employed the “Vanilla” model of Elman et al [13] as can be seen in Figure 2.2, in which the hidden steps are updated as follows:

$$h^t = f(W^i i^t + W^h h^{t-1} + b^h) \quad (2.20)$$

notation	meaning
$V$	Vocabulary size
$M$	word vector dimension
$n$	order of language model
$w$	a word; its index in the vocabulary
$H$	hidden layer size
$v_w$	one-hot vector for word $w$ . A vector with size $V$ with all null except the $w^{th}$ index (=1)
$X^T$	the matrix / vector $X$ transposed
$X * Y$	The element-wise matrix multiplication between matrices $X$ and $Y$
$i^t \in \mathbb{R}^M$	input vector of the network at time $t$ .
$h^t \in \mathbb{R}^V$	hidden layer of the network at time $t$
$o^t \in \mathbb{R}^V$	output layer of the network which expresses the unnormalised probability distribution at time $t$
$p^t \in \mathbb{R}^V$	the output vector of the network which is normalised for the conditional probability distribution
$f$	nonlinear function
$W^i \in \mathbb{R}^{H \times M}$	the weight matrix connecting the input layer with the hidden layer
$W^h \in \mathbb{R}^{H \times H}$	the weight matrix connecting the hidden layer of the previous step with the current hidden layer
$b^h \in \mathbb{R}^H$	the bias vector for the hidden layer
$W^o \in \mathbb{R}^{H \times V}$	the weight matrix connecting the hidden layer with the output layer
$b^o \in \mathbb{R}^V$	the bias vector for the output layer
$\mathcal{L}$	The objective (loss) function for learning the network
$dX / dx$	The Jacobian matrix - the matrix/vector of all of the first order partial derivatives of the loss function $\mathcal{L}$ with respect to the matrix $X$ or vector $x$ .

Table 2.3: Notations for recurrent neural network layers.

The activation function  $f$  can be either Tangent Hyperbolic, Sigmoid or ReLU as mentioned before. The starting state  $h_0$  is set to 0 to denote the initial state of the memory. In each time step, the RNNLM can optionally produce the probability distribution for a predicted word, given the sequence that network has scanned previously. The probability distribution over the vocabulary is derived similarly to the feed-forward networks:

$$\begin{aligned} o^t &= W^o h^t + b^o \\ p^t &= \text{softmax}(o^t) \end{aligned} \tag{2.21}$$

with the Softmax function explained in Equation 2.10. To be clear,  $o^t$  and  $p^t$  denote the unnormalised and normalised distribution generated at time step  $t$ . For the language modeling scenario, the input and output samples of the network in each training iteration are two sequences  $i$  and  $y$  in which the output sequence is the shift-by-1 version of the input sequence. The parameter set of the network including  $W^i, W^h, b^h, W^o$  and  $b^o$  are shared across time steps.

#### 2.2.4.1 Training Recurrent Networks

Similarly to the feed-forward models, The recurrent models can be efficiently trained with stochastic gradient descent (SGD). However, since the networks contain shared parameters at arbitrary numbers of time steps, the gradients are computed differently using back-propagation through time (BPTT). It can be observed that, a change in the parameters in an arbitrary time step  $t$  can lead to the change of the objective function in **all** subsequent steps.

The details of this algorithm are divided into two parts: the local gradients computed at each time step, and the global gradients accumulated by un-folding the networks.

**At each time step** The input and output sequences are denoted as  $i$  and  $y$ . At each time step  $t$ , the network receives one input  $i^t$  and predicts one output word  $y^t$ . Similar to the back-propagation process for feed-forward neural language models, we need to derive the derivatives of the per-time-step loss function  $\mathcal{L}^t$  with respect to the RNN parameters  $\{W^i, W^h, b^h, W^o \text{ and } b^o\}$  and the inputs  $\{i^t, h^{t-1}\}$ . Notably, the network at each time step receives two inputs: the input word vector  $i^t$  and the previous hidden layer  $h^{t-1}$ .

The Loss (Objective) function is identical to the feed-forward network:

$$\mathcal{L}^t = -\ln p^t \quad (2.22)$$

with  $p^t$  is the conditional probability of the output word at time  $t$  given the history from time 1 to  $t - 1$ , derived with Equation 2.21. The corresponding gradients are as follows:

$$do_j^t = \begin{cases} 1 - p_j^t & \text{if } j == y^t \\ -p_j^t & \text{otherwise} \end{cases} \quad (2.23)$$

The derivatives at the hidden layer can be derived:

$$\begin{aligned} dh^t &= W^{oT} do^t \\ dW^o &= do^t h^{tT} \\ db^o &= do^t \end{aligned} \quad (2.24)$$

After that, the errors are propagated to the input and the previous hidden layer. First of all, we use the simplified version of the RNN formulation in Equation 2.20, by denoting  $S = [W^i W^h]$  and  $z^t = [i^t; h^{t-1}]$ . The RNN formulation is neatly simplified as:

$$h^t = f(Sz^t + b^h) \quad (2.25)$$

We can derive the derivatives for the necessary weights ( $S$ ) and layers ( $z^t$ ) as follows:

$$\begin{aligned} dz^t &= S^T (f'(Sz^t + b^h) * dh^t) \\ dS &= (f'(Sz^t + b^h) * dh^t) z^{tT} \\ db^h &= d(Sz^t + b^h) \end{aligned} \quad (2.26)$$

From there, we can compute the gradients of the original variables ( $W^i, W^h$ ) and layers ( $i^t, h^{t-1}$ ).

$$\begin{aligned}
di^t &= W^{iT}(f'(Sz^t + b^h) * dh^t) \\
dh^{t-1} &= W^{hT}(f'(Sz^t + b^h) * dh^t) \\
dW^i &= (f'(Sz^t + b^h) * dh^t)i^{tT} \\
dW^h &= (f'(Sz^t + b^h) * dh^t)h^{t-1T} \\
db^i &= (f'(Sz^t + b^h) * dh^t)
\end{aligned} \tag{2.27}$$

The derivatives of the activation function has already been provided for the feed-forward neural network language model, in Equation 2.17.

**Back-propagation through time** The idea of BPTT is that the network is assumed to have different weights at each time-step. During training, the gradients are back-propagated to the previous layers through the recurrent connections. We compute the gradients for the weights at each time step with the formulation that we have just derived for a single time-step. Eventually, we force the shared weights at each time step to have the same value by **accumulating** the gradients together.

---

**Algorithm 1:** BPTT algorithm for “vanilla” RNNs

---

```

1 Inputs: input sequence  $i$  and output sequence  $y$  with length  $T$ .
2 Initialize the gradients  $do^t, dW^o, dW^i, dW^h, di^t, dh^{t-1}$  as zero vectors and matrices.
3 for  $t = T \rightarrow 1$  do
4     // At the output layer
5      $do^t \leftarrow v_{y_t} - p_t$ 
6      $dW^o \leftarrow dW^o + do^t \cdot h^{tT}$ 
7      $dh^t \leftarrow dh^t + W^{oT} do^t$ 
8     // RNN Backprop
9      $dW^i \leftarrow dW^i + (f'(Sz^t + b^h) * dh^t)i^{tT}$ 
10     $dW^h \leftarrow dW^h + (f'(Sz^t + b^h) * dh^t)h^{t-1T}$ 
11    // To the input layer and previous hidden layer
12     $di^t \leftarrow di^t + W^{iT}(f'(Sz^t + b^h) * dh^t)$ 
13     $dh^{t-1} \leftarrow dh^{t-1} + W^{hT}(f'(Sz^t + b^h) * dh^t)$ 
14 end

```

---

The whole BPTT process is illustrated in algorithm 1. The main difference between

BPTT and BP for feed-forward network is that, the gradients at the hidden state can be propagated by two paths: a “vertical” line from the hidden layer to the input layer, and the “horizontal” line to go back to the previous step. The network becomes harder to train when the sequence length  $T$  is too large, and the Vanilla RNN is not efficient to learn to remember long sequences.

**Problems with BPTT** Although truncated back-propagation through time provides a practical training method for RNNs, the nonlinear iterative nature of the simple RNN architecture still makes capturing long-term dependencies difficult. The two common problems encountered while training RNNs are the *exploding* and the *vanishing* gradients [2, 44]. On the one hand, the gradients can be exponentially large as in the back-propagation through time process which is detrimental for learning. On the other hand, we can also experience the phenomenon that gradients go quickly towards zero after being propagated through time steps. Consequently, the model is not able to track the signal and completely loses the memory trace in the past. For example, the original RNNLM normally has to truncate the BPTT at about 5 – 10 steps, which has the same modeling capacity and performance with 10-gram feed-forward models [20, 39]. The gradient exploding problem can be tackled adequately using gradient clipping. Pascanu et al [44] suggested to clip the norm of the gradients (for all parameters): given a gradient vector  $dx$  that is computed with BPTT, if the norm  $\|dx\|$  is greater than a threshold value  $\delta$ , then  $dx$  would be softly scaled:

$$dx \leftarrow dx \frac{\|dx\|}{\delta} \quad (2.28)$$

**Dealing with gradient vanishing** A number of solutions were proposed to solve the gradient vanishing problem. We will make a brief review of the most prominent approaches. First, Mikolov et al [40] proposed to integrate another memory layer which is formed by the bag-of-words addition of the input words over time, decays slower than the main hidden memory and is initialised as an identity matrix. The same initialisation trick is applied together with using Rectified Linear Units (ReLU) as the activation function in [33]. While both works mentioned are fairly simple, the RNNs can also be trained efficiently with second order derivatives using Hessian-Free optimization [37]. Even though the method was



prove to allow the network to acquire a reliable memory which can remain stable after hundreds of time steps, it is not easy to implement efficiently compared to traditional back-propagation. The most successful method that is applied to sequential modeling in general and language modeling in particular is the Long-Short Term Memory LSTM networks [24] that use an explicit memory cell combined with a gating mechanism to intensively deal with the gradient vanishing problem.

**LSTM Structure** The intuition of an LSTM starts from the integration of a linear memory unit, so that the gradient can flow smoothly during the back-propagation through time steps using a memory cell  $c_t$ .

$$c_t = c_{t-1} + f(Wx_i + Uhi - 1 + b)h_t = c_t \quad (2.29)$$

This approach is referred as “Leaky integration units” [5]. In the BPTT process, the gradient can flow over exactly one path through the memory units  $c_i$ , and since  $dc_i = dc_{i-1}$ , the gradients are guaranteed to not vanish. The recurrent architecture should also be able to be adequately robust to train long sequences, where there are certain inputs which are irrelevant to the modeling task. Sometimes, the memory of the network should be refreshed, for example at the beginning of a new utterance in Speech Recognition [17] or a new sentence in Machine Translation [56]. Hochreiter and Schmidhuber [24] enhanced the architecture by adding flexible and trainable gates that allows the RNN to reset the memory, control the amount of input and output respectively. The adaptive gates are built from the current input  $x_t$  and the previous hidden memory  $h_t$ .

The gates of the network include: the forget gate  $f_t$  is used to directly control the memory flow  $c_t$  to cut the connection with the previous steps, the input gate  $i_t$  decides the amount of input to be incorporated, the output gate  $o_t$  controls the amount of memory flow to be produced for the task and finally the candidate memory unit  $\tilde{C}$  that contributes to the current memory flow. All gates are defined similarly, with the first three gates use the Sigmoid activation to force the values to be in  $\{0, 1\}$ , while the candidate memory uses the Tanh activation function.

$$\begin{aligned}
f_t &= \text{Sigmoid}(W_f x_t + U_f h_t + b_f) \\
i_t &= \text{Sigmoid}(W_i x_t + U_i h_t + b_i) \\
o_t &= \text{Sigmoid}(W_o x_t + U_o h_t + b_o) \\
\tilde{C}_t &= \text{Sigmoid}(W_c x_t + U_c h_t + b_c)
\end{aligned} \tag{2.30}$$

In the next step, we decide the new information to be stored in the new memory cell. The cell is updated by combine the input gate and the candidate memory unit. Also, the forget gate is employed to drop certain information from the previous memory cell. Consequently, we come up with a new memory cell as follows:

$$C_t = f_t * C_t - 1 + i_t * \tilde{C}_t \tag{2.31}$$

Finally, we update the hidden state with the new cell state and the output gate:

$$h_t = o_t * \text{Tanh}(C_t) \tag{2.32}$$

The implementation of LSTM can be efficient by computing all gates in one single matrix multiplication, then applying the activation functions on different parts of the output. In practice, one can experience different implementation variations of LSTMs and RNNs in terms of initialisation, bias usage or different gate implementations such as the Gated Recurrent Unit [10]. The empirical research of [62] shows that there is not any substantial difference in terms of performance between different LSTM variations.

**Multi-layer recurrent neural network** We can extend a recurrent neural network by stacking the recurrent layers on top of each other in one single time-step. Concretely, the hidden layer at level  $i$  is the input of the hidden layer at level  $i + 1$ . For the recurrent connection of the network, the hidden layer at level  $i$  only depends on the the previous hidden layer of the same level  $i$ .

**Regularisation in RNN** Similar to feed-forward networks, the recurrent networks are prone to overfitting. In order to tackle the problem, it is efficient to apply dropout before the input of each recurrent unit [45, 63].

## 2.3 Convolutional Neural Networks

In this section, we provide the background information about the general Convolutional Neural Networks (CNN) which are mostly applied in the field of Computer Vision. We briefly cover the basic computation flow for the Convolutional layer - the fundamental component of the networks, whose intuitions are somewhat easier to understand for Computer Vision use case. The application of CNN for Natural Language Processing and Language Modeling in particular will be described in the next chapter.

### 2.3.1 The convolutional layer

Fundamentally, the convolutional layer is a biologically inspired variant of the fully connected layer (typical layers used in the feed-forward neural network language models). While the fully-connected layer expects a vector (1-dimension) as input and outputs a new vector using matrix multiplication, the convolutional layer receives inputs up to 3-dimension (such as images with channels, width and height) and, in most cases, produces an output that keeps the same dimensionality. Originally, the network was designed under the form of **Time-delayed neural networks** (TDNN) [60] in order to deal with 2D inputs such as speech signals. To clarify, the term TDNN and CNN are different names of a single concept that appeared in the literature at the same time, which refers to the use of *Local receptive field* in the input [16, 50].

The core idea of the layer is that, the neurons in the CNN output are only connected to a small region (local receptive field) of the layer before it, instead of all neurons in a fully-connected manner. The purpose of local receptive fields is to find features within the input that are invariant to position. For example, if we rotate an image then the textures are not changed, or if we look for the adjectives in the a sentence then the position of them is not important. The basic computation is to convolve (sliding) a window function applied to the input. The operation is illustrated on Figure 2.3. Each index in the output is obtained by performing the values of the kernel element-wise with the corresponding values in the input where the kernel is sliding on, then summing them up. The full convolution is done by repeating the operation by sliding the kernel over the input.




Figure 2.3: Simple illustration for convolution. The input is a 2D image, the output is obtained by sliding the kernel through the image.

**Forward propagation** For the purpose of simplicity, in the following operations, we derive the forward and backward formulas for 2D inputs and the delay between convolutional kernels is just one unit.

In order to mathematically formulate the convolutional layer, let us assume that the input is a  $N \times N$  square neuron layer  $X$  which is followed by the convolutional layer. Let the filter ( $K$ ) size be  $m \times m$ , the convolutional output  $Y$  will be of size  $(N - m + 1) \times (N - m + 1)$ . Each unit  $Y_{ij}$  in the output layer is computed by summing up the contributions from the input layer which are weighted by the kernel:

$$Y_{ij} = \sum_{a=1}^m \sum_{b=1}^m K_{ab} \cdot X_{(i+a)(j+b)} \quad (2.33)$$

The convolutional outputs are usually fed into non-linear functions in a similar manner to the fully-connected layers. In the convolutional layer, we specify the kernels as **learnable parameters** which are automatically adjusted during training based on the task of the network. For example, in the image classification tasks, CNNs have been found to be able to extract from low level features such as edges, colors, shapes to higher-level features such as facial shapes [32].

**Backward propagation** Let us assume that we have the loss function  $\mathcal{L}$  and we know the error values at the convolutional output,  $dY_{ij}$  (we keep the same notation as previous network types). Based on the output gradients, we want to compute the derivatives of the loss function with respect to the weights  $dK_{ij}$  ( $i$  and  $j$  are arbitrary indices).

First, using the chain rule, we sum all of the contributions of all expressions in which the variable occurs:

$$dK_{ab} = \sum_{i=1}^m \sum_{j=1}^m dY_{ij} \cdot X_{(i+a)(j+b)} \quad (2.34)$$

So that we can update the kernel parameters when we receive the back-propagation signal from the convolutional output. Furthermore, we also need to back-propagate the errors back to the convolutional input, in other words computing  $dX_{ij}$

$$dX_{ij} = \sum_{a=1}^m \sum_{b=1}^m K_{ab} \cdot dY_{(i-a+1)(j-b+1)} \quad (2.35)$$

From Equations 2.34,2.35, the backward operations can be efficiently implemented as convolutions similar to the forward operations. Also, the formulas also suggest that, we can expand the convolution with **padding** the inputs (so that the backward operations make sense for input units which are close to the border) and increasing the **stride** of the convolutions.

### 2.3.2 Hyper parameters for Convolutional Neural Network

Given the explanation of convolution - which is the connectivity between the neurons in the input layer and output layer, we can decide the number of neurons in the output layer (which has not been explicitly mentioned in the previous sections). The output layer size depends on the following hyper parameters (which are typically tried experimentally, or with a grid-search strategy).

- **Kernel size** The size of the window that we use to convolves the input.
- **Convolution depth** it corresponds to the number of kernels (sliding windows that we use to scan the input, each of which learns a different feature of the input. For

example, if we use convolutional neural networks to extract features for sentence classification, then multiple kernels can be distributed to learn features related to nouns, verbs or adjectives.

- As mentioned above, we can alter the step which we slide the filter over the input with the parameters **stride**. The larger the stride is, the smaller output is produced by convolution.
- We can also control the output size by adding zero values to the the borders of the input. For example in Figure 2.3, the convolution operator was not able to be performed at the index  $(1, 1)$  unless we pad 2 zero neurons to each dimension (and each side) of the input. In that case, we will obtain an output with the same size as the input.

Two hyper-parameters that affect the number of free parameters are kernel size and depth, while the other two control the output size of the layer.

## Chapter 3

# Convolutional Neural Language Models

The convolutional neural network have been successfully applied in Computer Vision [21, 32, 57] and Speech Recognition [1, 17]. The main idea of applying CNNs into those problems is to extract features related to short-time lags [24] - the arrangement of input units in a local area. For example, the CNNs are effective at extracting edges and tiny patterns in vision, or sequences of phonemes in speech signals.

A prominent work of applying CNNs in NLP is from Collobert et al. [11] who propose a neural network model with a convolutional layer in order to extract  $n$ -gram features for multiple NLP tasks including language modeling and part-of-speech tagging, chunking, named entity recognition. The network resembles the feed-forward neural language model [3], which includes the first layer that learns the continuous word representation and subsequent hidden layers. The main difference is that Collobert et al. [11] employ a convolutional layer (referred in the original work as Time-delayed neural network). Unfortunately, due to the computational limitations, the authors could not train the network by minimising the perplexity of the training data, but instead using a ranking-based loss function and focused on training the word embeddings. As mentioned before, the main purpose of the thesis is to investigate in the comparison between the standard feed-forward neural network and the convolutional neural network language model. Our CNN is constructed by extending a feed-forward language model (FFLM) with convolutional layers. In what follows, we first explain the implementation of the base FFLM and next, we describe the CNN that we study.

### 3.1 Base FFLM

Our baseline feed-forward language model (FFLM) is almost identical to the original model proposed by Bengio et al. [4], with only slight changes to push its performance as high as we can, producing a very strong baseline. In particular, we extend it with highway layers and use Dropout as regularization. The model is illustrated in Figure 3.1 and works as follows. First, each word in the input  $n$ -gram is mapped to a low-dimensional vector (viz. embedding) through a shared lookup table. Next, these word vectors are concatenated and fed to a highway layer [54].

**Highway layer** Highway layer is a variance of the conventional fully-connected layer. The mechanism of highway layers is to can combine the non-linear affine transformation of an input layer with itself to create the output hidden layer. Assuming the input and output of the highway layer is  $H_I$  and  $H_O$ . The materials for a highway layer is a typical non-linear transformation  $H$  and a transform gate  $T$

$$\begin{aligned} H &= f(H_I \cdot W_h + b_h) \\ T &= (H_I \cdot W_t + b_t) \end{aligned} \tag{3.1}$$

$W_h, b_h$  and  $W_t, b_t$  are linear connected weights for the model,  $f$  is a nonlinear function (ReLU in our work). Notably, the highway layer requires double the weights than a typical linear connected layer. The output of the highway layer is the combination of the transformed input and a part of the input carried away to the output.

$$H_O = H * T + H_I * (1 - T) \tag{3.2}$$

Succinctly, highway layers improve the gradient flow of the network by computing as output a convex combination between its input (called the *carry*) and a traditional non-linear transformation of it (called the *transform*). As a result, if there is a neuron whose gradient cannot flow through the transform component (e.g., because the activation is zero), it can still receive the back-propagation update signal through the carry gate. We empirically observed the usage of a single highway layer to improve importantly the performance of






Figure 3.1: Overview of baseline FFLM.

the model. Even though a systematic evaluation for this model is beyond the scope of the current paper, our empirical results demonstrate that it is a very competitive one (see Chapter 4).

Another important technique that we applied in our model is **Dropout**. Dropout [23] is a technique used to prevent overfitting

Finally, a softmax layer computes the model prediction for the upcoming word. We use ReLU for all non-linear activations and Dropout [23] is applied between each hidden layer.

## 3.2 CNN and variants

The proposed CNN network is produced by injecting a convolutional layer right after the words in the input are projected to their embeddings (Figure 3.2). Rather than being concatenated into a long vector, the embeddings  $x_i \in \mathbb{R}^k$  are concatenated transversally producing a matrix  $x_{1:n} \in \mathbb{R}^{n \times k}$ , where  $n$  is the size of the input and  $k$  is the embedding size. This matrix is fed to a time-delayed layer, which convolves a sliding window of  $w$  input vectors centered on each word vector using a parameter matrix  $W \in \mathbb{R}^{w \times k}$ . Convolution

is performed by taking the dot-product between the kernel matrix  $W$  and each sub-matrix  $x_{i-w/2:i+w/2}$  resulting in a scalar value for each position  $i$  in input context. This value represents how much the words encompassed by the window match the feature represented by the filter  $W$ . A ReLU activation function is applied subsequently so negative activations are discarded. This operation is repeated multiple times using various kernel matrices  $W$ , learning different features independently. Here we tie the number of learned kernels to be the same as the embedding dimensionality  $k$ . Thus, the output of this stage will be another matrix of dimensions  $n \times k$  containing the activations for each kernel at each time step.

Next, we add a batch normalization stage immediately after the convolutional output, which facilitates learning by addressing the internal covariate shift problem and regularizing the learned representations [26].

Finally, this feature matrix is directly fed into a fully connected layer that can project the extracted features into a lower-dimensional representation. This is different from previous work, where a max-over-time pooling operation was used to find the most activated feature in the time series. Our choice is motivated by the fact that the max pooling operator loses the specific position where the feature was detected, which is important for word prediction.

After this initial convolutional layer, the network proceeds identically to the FFNN by feeding the produced features into a highway layer, and then, to a softmax output.

This is our basic CNN architecture. We also experiment with three expansions to the basic model, as follows. First, we generalize the CNN by extending the shallow linear kernels with deeper multi-layer perceptrons, in what is called a Network-in-Network (**NIN**) structure [36]. This allows the network to produce non-linear filters, and it has achieved state-of-the-art performance in object recognition while reducing the number of total layers compared to other mainstream networks. Concretely, we implement NIN networks by using another convolutional layer with a  $1 \times 1$  kernel on top of the convolutional layer output. Notably, we do not include the global pooling layer in the original work.

Second, we explore stacking convolutional layers on top of each other (Multi-layer CNN or **ML-CNN**) to connect the local features into broader regional representations, as commonly done in computer vision. While this proved to be useful for sentence representation [27], here we have found it to be rather harmful for language modeling, as shown in Section ??.

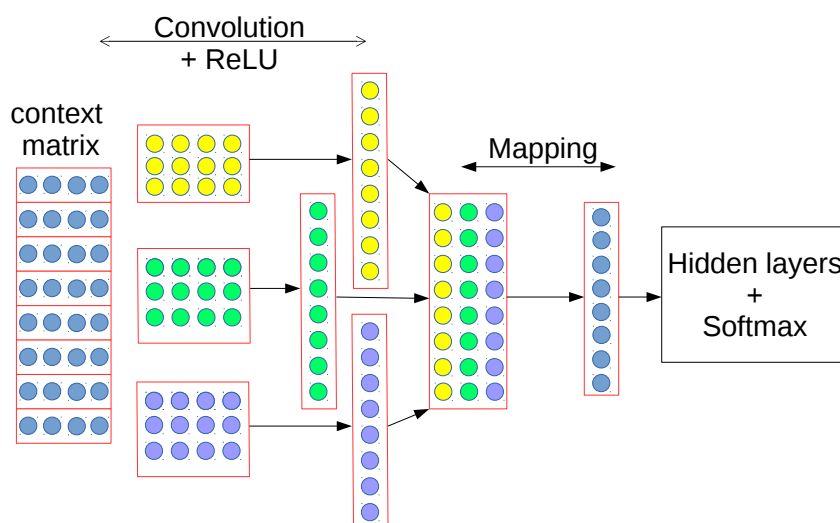


Figure 3.2: Convolutional layer on top of the context matrix.

Finally, we consider combining features learned through different kernel sizes (**COM**), as depicted in Figure 3.3. For example, we can have a combination of kernels that learn filters over 3-grams with others that learn over 5-grams. This is achieved simply by applying in parallel two or more sets of kernels to the input and concatenating their respective outputs [28].




Figure 3.3: Combining kernels with different sizes. We concatenate the outputs of 2 convolutional blocks with kernel size of 5 and 3 respectively.

# Chapter 4

## Experiments

### 4.1 Experiment details

We evaluate our model on two English corpora of different sizes and genres that have been used for language modeling evaluation before. The **Penn Treebank** contains one million words of newspaper text with 10K words in the vocabulary. We reuse the preprocessing and train/test/valid division from [40]. **Europarl-NC** is a 64-million word corpus that was developed for a Machine Translation shared task [6], combining Europarl data (from parliamentary debates in the European Union) and News Commentary data. We preprocessed the corpus with tokenization and true-casing tools from the Moses toolkit [31]. The vocabulary is composed of words that occur at least 3 times in the training set and contains approximately 60K words.

We train our models using Stochastic Gradient Descent (SGD) and we reduce the learning rate by a fixed proportion every time the validation perplexity increases after one epoch. The values for learning rate, learning rate shrinking and mini-batch sizes as well as context size are fixed once and for all based on insights drawn from previous work [12, 19, 55] and through experimentation with the Penn Treebank validation set. Experimentally we found SGD to be efficient and adequately fast, while other learning algorithms involve additional hyper parameters (such as alpha in RMSprop [58])

Specifically, the learning rate is set to 0.05, with mini-batch size of 128 (we do not take the average of loss over the batch, and the training set is shuffled). We multiply the learning

rate by 0.5 every time we shrink it and clip the gradients if their norm is larger than 12. The network parameters are initialized randomly on a range from  $-0.01$  to  $0.01$  and the context size is set to 16. In Section ?? we show that this large context window is fully exploited.

For the base FFNN and CNN we study embedding sizes (and thus, number of kernels)  $k = 128, 256$ . For  $k = 128$  we explore the simple CNN, incrementally adding NIN and COM variations (in that order) and, alternatively, using a ML-CNN. For  $k = 256$ , we only explore the former three alternatives. For the kernel size, we set it to  $w = 3$  words for the simple CNN (out of options 3, 5, 7, 9), whereas for the COM variant we use  $w = 3$  and 5. However, we observed the models to be generally robust to this parameter. Dropout rates are tuned specifically for each combination of model and dataset based on the validation perplexity. We also add small dropout ( $0.05 - 0.15$ ) when we train the networks on the small corpus (Penn Treebank).

The experimental results for recurrent neural network language models, such as Recurrent Neural Networks (RNN) and Long-Short Term Memory (LSTM), on the Penn Treebank are quoted from previous work; for Europarl-NC, we train our own models (we also report the performance of these in-house trained RNN and LSTM models on the Penn Treebank for reference). Specifically, we train LSTMs with embedding size  $k = 256$  and number of layers  $L = 2$  as well as  $k = 512$  with  $L = 1, 2$ . We train one RNN with  $k = 512$  and  $L = 2$ . To train these models, we use the published source code from [63]. Our own models are also implemented in Torch<sup>7</sup> for easier comparison.

For all models trained on Europarl-NC, we used hierarchical softmax to speed up training, clustering the words in the vocabulary into 245 classes based on word frequency.

## 4.2 Results

Our experimental results are summarized in Table 4.1.

First of all, we can see that even though the FFNN gives a very competitive performance, the addition of convolutional layers is clearly effective to increase it even further<sup>2</sup>.

---

<sup>7</sup>We'll link the source code here if the paper is accepted

<sup>2</sup>In our experiments, increasing the number of fully connected layers is harmful. Two hidden layers with highway connections is the best setting we could find.

Concretely, we observe a solid 13% reduction of perplexity compared to the feed-forward network after using Network-in-Network in all setups for both corpora. CNN alone yields a 6% improvement, while MLPConv, in line with our expectations, adds another 5% reduction in perplexity. A final (smaller) improvement comes from combining kernels of size 3 and 5, which can be attributed to a more expressive model that can learn patterns of  $n$ -grams of different sizes. In contrast to the successful two variants above, the multi-layer CNN did not help in better capturing the regularities of text, but rather the opposite: the more convolutional layers were stacked, the worse the performance. This also stands in contrast to the tradition of convolutional networks in Computer Vision, where using very deep convolutional neural networks is key to having better models. In text, deep convolution for text representation is rather rare, and has only been applied in sentence representation [27]. We conjecture that the reason why deep CNNs may not be so effective for text could be the non-recursive nature of the textual data after convolution. In other words, the convolution output for an image can be construed to be a new image, which yet again can be subject to new convolution operations, whereas the textual counterpart may no longer have the same property.

Regarding the comparison with a stronger LSTM, our models can perform competitively under the same embedding dimension. However, the LSTM can be easily scaled using larger models, as shown in Zaremba et al. [63], which gives the best known results to date. This is not an option for our model which heavily overfits with large hidden layers (around 1000) even with very large dropout values.

To sum up, we have established that the results of our CNN model are well above those of simple feed forward networks and recurrent neural networks. While they are below state of the art LSTMs, they are able to perform competitively with them for small and moderate-size models. Scaling to larger sizes may be today the main roadblock for CNNs to reach the same performances as large LSTMs in language modeling.

Model	$k$	ker	Penn Treebank			Europarl-NC		
			val	test	#p	val	test	#p
Vanilla FFNN	128	-	156	147	4.5M	-	-	-
Baseline FFNN	128	-	114	109	4.5M	-	-	-
+CNN	128	3	108	102	4.5M	-	-	-
+MLPConv	128	3	102	97	4.5M	-	-	-
+MLPConv+COM	128	3+5	96	<b>92</b>	8M	-	-	-
+ML-CNN (2 layers)	128	3	113	108	8M	-	-	-
+ML-CNN (4 layers)	128	3	130	124	8M	-	-	-
Vanilla FFNN	256	-	161	152	8.2M	-	-	-
Baseline FFNN	256	-	110	105	8.2M	133	174	48M
+CNN	256	3	104	98	8.3M	112	133	48M
+MLPConv	256	3	97	93	8.3M	107	128	48M
+MLPConv+COM	256	3+5	95	<b>91</b>	18M	108	<b>128</b>	83M
+MLPConv+COM	512	3+5	96	92	52M	-	-	-
Model	hid	#l	Penn Treebank			Europarl-NC		
			val	test	#p	val	test	#p
RNN [40]	300	1	133	129	6M	-	-	
LSTM [40]	300	1	120	115	6.3M	-	-	
LSTM [63]	1500	2	82	<b>78</b>	48M	-	-	
LSTM (trained in-house)	256	2	108	103	5.1M	137	155	31M
LSTM (trained in-house)	512	1	123	118	12M	133	149	62M
LSTM (trained in-house)	512	2	94	90	11M	114	<b>124</b>	63M
RNN (trained in-house)	512	2	129	121	10M	152	173	61M

Table 4.1: Results on Penn Treebank and Europarl-NC. Figure of merit is perplexity (lower is better). Legend:  $k$ : embedding size;  $ker$ : kernel size;  $val$ : results on validation data;  $test$ : results on test data;  $\#p$ : number of parameters;  $hid$ : size of hidden layers;  $\#l$ : number of layers.



no matter how are afraid how question is how remaining are how to say how	as little as of more than as high as as much as as low as	a merc spokesman a company spokesman a boeing spokesman a fidelity spokesman a quotron spokeswoman	amr chairman robert chief economist john chicago investor william exchange chairman john texas billionaire robert
would allow the does allow the still expect ford warrant allows the funds allow investors	more evident among a dispute among bargain-hunting among growing fear among paintings listed among	facilities will substantially which would substantially dean witter actually we 'll probably you should really	have until nov. operation since aug. quarter ended sept. terrible tuesday oct. even before june

Figure 5.1: Some example phrases that have highest activations for 8 example kernels (each box), extracted from the validation set of the Penn Treebank. Model trained with 256 kernels for 256-dimension word vectors.

## Chapter 5

### Model Analysis

In what follows, we obtain insights into the inner workings of the CNN by looking into the linguistic patterns that the kernels learn to extract and also studying the temporal information extracted by the network in relation to its prediction capacity.

**Learned patterns** To get some insight into the kind of patterns that each kernel is learning to detect, we fed trigrams from the validation set of the Penn Treebank to each of the kernels, and extracted the ones that most highly activated the kernel. Some examples are shown in Figure 5.1. Since the word windows are made of embeddings, we can expect




Figure 5.2: The distribution of positive weights over context positions, where 1 is the position closest to the predicted word.

patterns with similar embeddings to have close activation outputs. This is borne out in the analysis: The kernels specialize in distinct features of the data, including more syntactic-semantic constructions (cf. the “comparative kernel” including *as ... as* patterns, but also *of more than*) and more lexical or topical features (cf. the “ending-in-month-name” kernel). Even in the more lexicalized features, however, we see linguistic regularities at different levels being condensed in a single kernel: For instance, the “spokesman” kernel detects phrases consisting of an indefinite determiner, a company name (or the word *company* itself) and the word “spokesman”. We hypothesize that the convolutional layer adds an “I identify one specific feature, but at a high level of abstraction” dimension to a feed-forward neural network, similarly to what has been observed in image classification [32].

**Temporal information** To the best of our knowledge, the longest context used in feed-forward language models is 10 [20], where no significant change in terms of perplexity was observed for bigger context sizes, even though in that work only same-sentence contexts were considered. In our experiments, we use a larger context size of 16 while removing the sentence boundary limit (as commonly done in  $n$ -gram language models) such that the network can take into account the words in the previous sentences.

To analyze whether all this information was effectively used, we took our best model, the CNN-NIN-COM model with embedding size of 256 (fourth line, second block in Table 4.1), and we identified the weights in the model that map the convolutional output (of size  $n \times k$ ) to a lower dimensional vector (the “mapping” layer in Figure 3.2). Recall that the output of the convolutional layer is a matrix indexed by time step and kernel index representing the activation of the kernel when convolved with a window of text centered around the given time step. Thus, output units of the above mentioned mapping predicate over an ensemble of kernel activations for each time-step. We can identify the patterns that they learn to detect by extracting the time-kernel combinations for which they have positive weights (since we have ReLU activations, negative weights are equivalent to ignoring a feature). First, we asked ourselves whether these units tend to be more focused on the time-steps closer to the target or not. To test this, we calculated the sum of the positive weights for each position in time using an average of the mappings that correspond to each output unit. The results are shown in Figure 5.2. As could be expected, positions that are close to the token to be predicted have many active units (local context is very informative; see positions 2-4). However, surprisingly, positions that are actually far from the target are also quite active (see positions 10-14, with a spike at 11). It seems like the CNN is putting quite a lot of effort on characterizing long-range dependencies.

Next, we checked that the information extracted from the positions that are far in the past are actually used for prediction. To measure this, we artificially lesioned the network so it would only read the features from a given range of timesteps (words in the context). To lesion the network we manually masked the weights of the mapping that focus on times outside of the target range by setting them to zero. We started using only the word closest to the final position and sequentially unmasked earlier positions until the full context was




Figure 5.3: Perplexity change over position, by incrementally revealing the Mapping’s weights corresponding to each position.

used again. The result of this experiment is presented in Figure 5.3, and it confirms our previous observation that positions that are the farthest away contribute to the predictions of the model. The perplexity drops dramatically as the first positions are unmasked, and then decreases more slowly. However, the last few positions still contribute to the final perplexity. Indeed, the decrease almost follows a power-law distribution, in which the perplexity decreases whenever we unmask a new position according to the following equation.

$$f(x) = 10e3 \times x^{-0.9} \quad (5.1)$$

# Chapter 6

## Related Works

Convolutional Neural Networks (CNNs) were originally designed to deal with hierarchical representation in Computer Vision [35]. Deep convolutional networks have been successfully applied in image classification and understanding [21, 52]. In such systems the convolutional kernels learn to detect visual features at both local and more abstract levels.

In NLP, CNNs have been mainly applied to static classification task for discovering latent structures in text. Kim [28] use a CNN to tackle sentence classification, with competitive results. This work also introduces kernels with varying window sizes to learn complementary features at different aggregation levels. Kalchbrenner et al. [27] propose a convolutional architecture for sentence representation that vertically stacks multiple convolution layers, each of which can learn independent convolution kernels. CNNs with similar structures have also been applied to other classification tasks, such as semantic matching [25], relation extraction [42] and information retrieval [51]. In contrast, Collobert et al. [11] explore a CNN architecture to solve various sequential and non-sequential NLP tasks such as part-of-speech tagging, named entity recognition and also language modeling. This is perhaps the work that is closest to ours in the existing literature. However, their model differs from ours in that it uses a max-pooling layer that picks the most activated feature across time, thus ignoring temporal information, whereas we explicitly avoid doing so. More importantly, the language models trained in this work are only evaluated through downstream tasks and through the quality of the learned word embeddings, but not on the sequence prediction task itself.

In this work we focus on statistical language modeling, which differs from most of the tasks where CNNs have been applied before in multiple ways. First, the input normally consists of incomplete sequences of words rather than complete sentences. Second, as a classification problem, it features an extremely large number of classes (the words in a large vocabulary). Finally, temporal information, which can be safely discarded in many settings with little impact in performance, is critical here: An  $n$ -gram appearing close to the predicted word may be more informative, or yield different information, than the same  $n$ -gram appearing several tokens earlier.

# Chapter 7

## Conclusion

In this work, we have investigated the use of Convolutional Neural Networks for language modeling, a sequential prediction task. We incorporate a CNN layer on top of a strong feed-forward model enhanced with modern techniques like Highway Layers and Dropout. Our results show a solid 13% improvement in perplexity with respect to the feed-forward model across two corpora of different sizes and genres when the model uses Network-in-Network and combine kernels of different window sizes. However, even without these additions we show CNNs to effectively learn language patterns to significantly decrease the model perplexity.

In our view, this improvement responds to two key properties of CNNs, highlighted in the analysis. First, as we have shown, they are able to integrate information from larger context windows, using information from words that are as far as 16 positions away from the predicted word. Second, as we have qualitatively shown, the kernels learn to detect specific patterns at a high level of abstraction. This is analogous to the role of convolutions in Computer Vision. The analogy, however, has limits; for instance, a deeper model stacking convolution layers harms performance in language modeling, while it greatly helps in Computer Vision. We conjecture that this is due to the differences in the nature of visual vs. linguistic data. The convolution creates sort of abstract images that still retain significant properties of images. When applied to language, it detects important textual features but distorts the input, such that it is not text anymore.

As for recurrent models, even if our model outperforms RNNs, it is well below state-of-the-art LSTMs. Since CNNs are quite different in nature, we believe that a fruitful line of future research could focus on integrating the convolutional layer into a recurrent structure for language modeling, as well as other sequential problems, perhaps capturing the best of both worlds.



# **Appendix A**

## **Publications by Author**

- 2016

–

- 2015

–

–

# Bibliography

- [1] Ossama Abdel-Hamid, Abdel-rahman Mohamed, Hui Jiang, and Gerald Penn. Applying convolutional neural networks concepts to hybrid nn-hmm model for speech recognition. In *2012 IEEE international conference on Acoustics, speech and signal processing (ICASSP)*, pages 4277–4280. IEEE, 2012.
- [2] Yoshua Bengio, Patrice Simard, and Paolo Frasconi. Learning long-term dependencies with gradient descent is difficult. *IEEE transactions on neural networks*, 5(2): 157–166, 1994.
- [3] Yoshua Bengio, Réjean Ducharme, Pascal Vincent, and Christian Jauvin. A neural probabilistic language model. *journal of machine learning research*, 3(Feb):1137–1155, 2003.
- [4] Yoshua Bengio, Holger Schwenk, Jean-Sébastien Senécal, Frédéric Morin, and Jean-Luc Gauvain. Neural probabilistic language models. In *Innovations in Machine Learning*, pages 137–186. Springer, 2006.
- [5] Yoshua Bengio, Nicolas Boulanger-Lewandowski, and Razvan Pascanu. Advances in optimizing recurrent networks. In *2013 IEEE International Conference on Acoustics, Speech and Signal Processing*, pages 8624–8628. IEEE, 2013.
- [6] Ondřej Bojar, Rajen Chatterjee, Christian Federmann, Barry Haddow, Matthias Huck, Chris Hokamp, Philipp Koehn, Varvara Logacheva, Christof Monz, Matteo Negri, Matt Post, Carolina Scarton, Lucia Specia, and Marco Turchi. Findings of the 2015 workshop on statistical machine translation. In *Proceedings of the Tenth Workshop on Statistical Machine Translation*, pages 1–46, Lisbon, Portugal, September

2015. Association for Computational Linguistics. URL <http://aclweb.org/anthology/W15-3001>.
- [7] Peter F Brown, Peter V Desouza, Robert L Mercer, Vincent J Della Pietra, and Jenifer C Lai. Class-based n-gram models of natural language. *Computational linguistics*, 18(4):467–479, 1992.
- [8] Ciprian Chelba and Frederick Jelinek. Structured language modeling. *Computer Speech & Language*, 14(4):283–332, 2000.
- [9] Stanley F Chen and Joshua Goodman. An empirical study of smoothing techniques for language modeling. *Computer Speech & Language*, 13(4):359–394, 1999.
- [10] Kyunghyun Cho, Bart Van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. Learning phrase representations using rnn encoder-decoder for statistical machine translation. *arXiv preprint arXiv:1406.1078*, 2014.
- [11] Ronan Collobert, Jason Weston, Léon Bottou, Michael Karlen, Koray Kavukcuoglu, and Pavel Kuksa. Natural language processing (almost) from scratch. *The Journal of Machine Learning Research*, 12:2493–2537, 2011.
- [12] Jacob Devlin, Rabih Zbib, Zhongqiang Huang, Thomas Lamar, Richard M Schwartz, and John Makhoul. Fast and robust neural network joint models for statistical machine translation. In *ACL (1)*, pages 1370–1380. Citeseer, 2014.
- [13] Jeffrey L Elman. Finding structure in time. *Cognitive science*, 14(2):179–211, 1990.
- [14] Marcello Federico, Nicola Bertoldi, and Mauro Cettolo. Irs1lm: an open source toolkit for handling large scale language models. In *Interspeech*, pages 1618–1621, 2008.
- [15] Denis Filimonov and Mary Harper. A joint language model with fine-grain syntactic tags. In *Proceedings of the 2009 Conference on Empirical Methods in Natural Language Processing: Volume 3-Volume 3*, pages 1114–1123. Association for Computational Linguistics, 2009.

- [16] Kunihiro Fukushima. Neocognitron: A self-organizing neural network model for a mechanism of pattern recognition unaffected by shift in position. *Biological cybernetics*, 36(4):193–202, 1980.
- [17] Alex Graves and Jürgen Schmidhuber. Framewise phoneme classification with bidirectional lstm and other neural network architectures. *Neural Networks*, 18(5):602–610, 2005.
- [18] Jiuxiang Gu, Zhenhua Wang, Jason Kuen, Lianyang Ma, Amir Shahroudy, Bing Shuai, Ting Liu, Xingxing Wang, and Gang Wang. Recent advances in convolutional neural networks. *CoRR*, abs/1512.07108, 2015.
- [19] Le Hai Son, Ilya Oparin, Alexandre Allauzen, Jean-Luc Gauvain, and François Yvon. Structured output layer neural network language model. In *Acoustics, Speech and Signal Processing (ICASSP), 2011 IEEE International Conference on*, pages 5524–5527. IEEE, 2011.
- [20] Le Hai Son, Alexandre Allauzen, and François Yvon. Measuring the influence of long range dependencies with neural network language models. In *Proceedings of the NAACL-HLT 2012 Workshop: Will We Ever Really Replace the N-gram Model? On the Future of Language Modeling for HLT*, pages 1–10. Association for Computational Linguistics, 2012.
- [21] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. *arXiv preprint arXiv:1512.03385*, 2015.
- [22] Kenneth Heafield. Kenlm: Faster and smaller language model queries. In *Proceedings of the Sixth Workshop on Statistical Machine Translation*, pages 187–197. Association for Computational Linguistics, 2011.
- [23] Geoffrey E. Hinton, Nitish Srivastava, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov. Improving neural networks by preventing co-adaptation of feature detectors. *CoRR*, abs/1207.0580, 2012. URL <http://arxiv.org/abs/1207.0580>.

- [24] Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural computation*, 9(8):1735–1780, 1997.
- [25] Baotian Hu, Zhengdong Lu, Hang Li, and Qingcai Chen. Convolutional neural network architectures for matching natural language sentences. In *Advances in Neural Information Processing Systems*, pages 2042–2050, 2014.
- [26] Sergey Ioffe and Christian Szegedy. Batch normalization: Accelerating deep network training by reducing internal covariate shift. *arXiv preprint arXiv:1502.03167*, 2015.
- [27] Nal Kalchbrenner, Edward Grefenstette, and Phil Blunsom. A convolutional neural network for modelling sentences. In *Proceedings of the 52nd Annual Meeting of the Association for Computational Linguistics, ACL 2014, June 22-27, 2014, Baltimore, MD, USA, Volume 1: Long Papers*, pages 655–665, 2014. URL <http://aclweb.org/anthology/P/P14/P14-1062.pdf>.
- [28] Yoon Kim. Convolutional neural networks for sentence classification. *arXiv preprint arXiv:1408.5882*, 2014.
- [29] Zipf George Kingsley. Selective studies and the principle of relative frequency in language, 1932.
- [30] Reinhard Kneser and Hermann Ney. Improved backing-off for m-gram language modeling. In *Acoustics, Speech, and Signal Processing, 1995. ICASSP-95., 1995 International Conference on*, volume 1, pages 181–184. IEEE, 1995.
- [31] Philipp Koehn, Hieu Hoang, Alexandra Birch, Chris Callison-Burch, Marcello Federico, Nicola Bertoldi, Brooke Cowan, Wade Shen, Christine Moran, Richard Zens, et al. Moses: Open source toolkit for statistical machine translation. In *Proceedings of the 45th annual meeting of the ACL on interactive poster and demonstration sessions*, pages 177–180. Association for Computational Linguistics, 2007.
- [32] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E Hinton. Imagenet classification with deep convolutional neural networks. In *Advances in neural information processing systems*, pages 1097–1105, 2012.

- [33] Quoc V Le, Navdeep Jaitly, and Geoffrey E Hinton. A simple way to initialize recurrent networks of rectified linear units. *arXiv preprint arXiv:1504.00941*, 2015.
- [34] B Boser Le Cun, John S Denker, D Henderson, Richard E Howard, W Hubbard, and Lawrence D Jackel. Handwritten digit recognition with a back-propagation network. In *Advances in neural information processing systems*. Citeseer, 1990.
- [35] Yann LeCun and Yoshua Bengio. Convolutional networks for images, speech, and time series. *The handbook of brain theory and neural networks*, 3361(10):1995, 1995.
- [36] Min Lin, Qiang Chen, and Shuicheng Yan. Network in network. *arXiv preprint arXiv:1312.4400*, 2013.
- [37] James Martens and Ilya Sutskever. Learning recurrent neural networks with hessian-free optimization. In *Proceedings of the 28th International Conference on Machine Learning (ICML-11)*, pages 1033–1040, 2011.
- [38] Tomas Mikolov, Martin Karafiát, Lukas Burget, Jan Cernocký, and Sanjeev Khudanpur. Recurrent neural network based language model. In *INTERSPEECH*, volume 2, page 3, 2010.
- [39] Tomáš Mikolov, Stefan Kombrink, Lukáš Burget, Jan Honza Černocký, and Sanjeev Khudanpur. Extensions of recurrent neural network language model. In *Acoustics, Speech and Signal Processing (ICASSP), 2011 IEEE International Conference on*, pages 5528–5531. IEEE, 2011.
- [40] Tomas Mikolov, Armand Joulin, Sumit Chopra, Michael Mathieu, and Marc’Aurelio Ranzato. Learning longer memory in recurrent neural networks. *arXiv preprint arXiv:1412.7753*, 2014.
- [41] Hermann Ney, Ute Essen, and Reinhard Kneser. On structuring probabilistic dependencies in stochastic language modelling. *Computer Speech & Language*, 8(1):1–38, 1994.
- [42] Thien Huu Nguyen and Ralph Grishman. Relation extraction: Perspective from convolutional neural networks. In *Proceedings of NAACL-HLT*, pages 39–48, 2015.

- [43] Thomas R Niesler and Philip C Woodland. A variable-length category-based n-gram language model. In *Acoustics, Speech, and Signal Processing, 1996. ICASSP-96. Conference Proceedings., 1996 IEEE International Conference on*, volume 1, pages 164–167. IEEE, 1996.
- [44] Razvan Pascanu, Tomas Mikolov, and Yoshua Bengio. On the difficulty of training recurrent neural networks. *ICML (3)*, 28:1310–1318, 2013.
- [45] Vu Pham, Théodore Bluche, Christopher Kermorvant, and Jérôme Louradour. Dropout improves recurrent neural networks for handwriting recognition. In *Frontiers in Handwriting Recognition (ICFHR), 2014 14th International Conference on*, pages 285–290. IEEE, 2014.
- [46] Ronald Rosenfeld. Two decades of statistical language modeling: Where do we go from here. In *Proceedings of the IEEE*, page 2000, 2000.
- [47] David E Rumelhart, Geoffrey E Hinton, and Ronald J Williams. Learning internal representations by error propagation. Technical report, DTIC Document, 1985.
- [48] Ivan A Sag, Timothy Baldwin, Francis Bond, Ann Copestake, and Dan Flickinger. Multiword expressions: A pain in the neck for nlp. In *International Conference on Intelligent Text Processing and Computational Linguistics*, pages 1–15. Springer, 2002.
- [49] Holger Schwenk. Continuous space language models. *Computer Speech & Language*, 21(3):492–518, 2007.
- [50] Thomas Serre, Lior Wolf, Stanley Bileschi, Maximilian Riesenhuber, and Tomaso Poggio. Robust object recognition with cortex-like mechanisms. *IEEE transactions on pattern analysis and machine intelligence*, 29(3):411–426, 2007.
- [51] Yelong Shen, Xiaodong He, Jianfeng Gao, Li Deng, and Grégoire Mesnil. A latent semantic model with convolutional-pooling structure for information retrieval. In *Proceedings of the 23rd ACM International Conference on Conference on Information and Knowledge Management*, pages 101–110. ACM, 2014.

- [52] Karen Simonyan and Andrew Zisserman. Very deep convolutional networks for large-scale image recognition. *arXiv preprint arXiv:1409.1556*, 2014.
- [53] Nitish Srivastava, Geoffrey E Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov. Dropout: a simple way to prevent neural networks from overfitting. *Journal of Machine Learning Research*, 15(1):1929–1958, 2014.
- [54] Rupesh Kumar Srivastava, Klaus Greff, and Jürgen Schmidhuber. Highway networks. *CoRR*, abs/1505.00387, 2015. URL <http://arxiv.org/abs/1505.00387>.
- [55] Sainbayar Sukhbaatar, Jason Weston, Rob Fergus, et al. End-to-end memory networks. In *Advances in Neural Information Processing Systems*, pages 2431–2439, 2015.
- [56] Ilya Sutskever, Oriol Vinyals, and Quoc V Le. Sequence to sequence learning with neural networks. In *Advances in neural information processing systems*, pages 3104–3112, 2014.
- [57] Christian Szegedy, Wei Liu, Yangqing Jia, Pierre Sermanet, Scott Reed, Dragomir Anguelov, Dumitru Erhan, Vincent Vanhoucke, and Andrew Rabinovich. Going deeper with convolutions. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 1–9, 2015.
- [58] Tijmen Tieleman and Geoffrey Hinton. Lecture 6.5-rmsprop: Divide the gradient by a running average of its recent magnitude. *COURSERA: Neural Networks for Machine Learning*, 4(2), 2012.
- [59] Ashish Vaswani, Yinggong Zhao, Victoria Fossum, and David Chiang. Decoding with large-scale neural language models improves translation. In *EMNLP*, pages 1387–1392. Citeseer, 2013.
- [60] Alexander Waibel, Toshiyuki Hanazawa, Geoffrey Hinton, Kiyohiro Shikano, and Kevin J Lang. Phoneme recognition using time-delay neural networks. *Acoustics, Speech and Signal Processing, IEEE Transactions on*, 37(3):328–339, 1989.



- [61] Ian H Witten and Timothy C Bell. The zero-frequency problem: Estimating the probabilities of novel events in adaptive text compression. *Ieee transactions on information theory*, 37(4):1085–1094, 1991.
- [62] Wojciech Zaremba. An empirical exploration of recurrent network architectures. 2015.
- [63] Wojciech Zaremba, Ilya Sutskever, and Oriol Vinyals. Recurrent neural network regularization. *arXiv preprint arXiv:1409.2329*, 2014.
- [64] Matthew D Zeiler. Adadelta: an adaptive learning rate method. *arXiv preprint arXiv:1212.5701*, 2012.